



Indian Institute of Technology Delhi

ELL 783 - Operating Systems

Assignment 2

Real Time Scheduling

Name: Mohammad Yusuf, Madhur Sharma

Entry No.: 2023EET2757, 2023EET2193

Course Instructor: Prof. Smruti Ranjan Sarangi

Introduction

This report details the implementation methodology of real-time scheduling policies, specifically the Earliest Deadline First (EDF) and Rate-Monotonic (RM) scheduling algorithms, within the xv6 operating system. The report includes code snippets and pseudocode to demonstrate the implementation approach.

EDF Scheduling Algorithm

The EDF (Earliest Deadline First) scheduling algorithm is a dynamic priority scheduling algorithm used in real-time operating systems to schedule processes based on their deadlines. In this report, we discuss the implementation of the EDF scheduling algorithm in the xv6 operating system.

Schedulability Test

In our implementation, we perform a schedulability test to determine whether a process is schedulable under the EDF policy. The utility of all scheduled processes is added together, and if the total utility exceeds a certain threshold, the process is exited and not considered for future scheduling checks. To ensure accuracy, we perform calculations using float `cpu_utilz` variable which is precise up to 4 decimal places. Below is a snippet of the code illustrating the schedulability test:

```
752     // EDF
753     if (policy == 0)
754     {
755         for (struct proc *p1 = ptable.proc; p1 < &ptable.proc[NPROC]; p1++){
756             if (p1->sched_policy == 0 || p1->pid == p->pid){
757                 //calculating utilization of all sched_policy=0 including the new process
758                 cpu_utilz += (float)(p1->exec_time) / p1->deadline;
759             }
760
761             if(cpu_utilz > 1){
762                 release(&ptable.lock);
763                 kill(pid);
764                 return -22;
765             }
766         }
767     }
```

Scheduling Algorithm

The EDF scheduling algorithm prioritizes processes based on their deadlines, ensuring that processes with earlier deadlines are scheduled first. We calculate the time left before the deadline for each runnable process and schedule the one with the least time left. In cases where multiple processes have the same time left, we prioritize the process with the lower PID. Here is a simplified version of the scheduling algorithm code:

Code

```
386     if (sch_policy == 0) //EDF
387     {
388         int closest_deadline = 100000000;
389         struct proc *p_selected = NULL;
390
391         for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
392         {
393             if (p->state == RUNNABLE && p->sched_policy == 0) {
394                 int p_time_left = (p->arrival_time + p->deadline) - ticks;
395
396                 if (p_time_left < closest_deadline){ //find earliest deadline
397                     closest_deadline = p_time_left;
398                     p_selected = p;
399                 }
400                 else if (p_time_left == closest_deadline){
401                     if (p->pid < p_selected->pid) //select lowest pid
402                         p_selected=p;
403                 }
404             }
405         }
406
407         if(p_selected!=NULL){
408             p=p_selected;
409
410             // Switch to chosen process. It is the process's job
411             // to release ptable.lock and then reacquire it
412             // before jumping back to us.
413             c->proc = p;
414             switchuvm(p);
415             p->state = RUNNING;
416
417             switch(&(c->scheduler), p->context);
418             switchkvm();
419             p->elapsed_time++;
420
421             // Process is done running for now.
422             // It should have changed its p->state before coming back.
423             c->proc = 0;
424         }
425     }
```

In conclusion, our implementation of the EDF scheduling algorithm in xv6 ensures that processes are scheduled based on their deadlines, allowing for efficient utilization of resources in a real-time operating system environment. The schedulability test and scheduling algorithm work together to prioritize processes

effectively, contributing to the overall performance and responsiveness of the system.

RMS Scheduling Algorithm

The RMS (Rate-Monotonic Scheduling) algorithm is a fixed-priority scheduling algorithm used in real-time operating systems to schedule processes based on their rates or priorities. This report presents the implementation and functioning of the RMS scheduling algorithm within the xv6 operating system environment.

Schedulability Test

Our implementation includes a schedulability test to determine whether a process is schedulable under the RMS policy. The test involves calculating the total utility of all scheduled processes and comparing it with precomputed `rms_bound[]` values, which is the Liu Layland Bound, based on the number of processes. If the total utility exceeds the `rms_bound[i]`, the process is terminated and not considered for future scheduling checks. The accuracy of calculations is maintained by using float arithmetic and `rms_bound[]` values are out of 100, which is precise up to 4 decimal points. Here is a snippet of the code illustrating the schedulability test:

```
769 // RMS
770 else if (policy == 1)
771 {
772     int num_rms_proc = 0;
773
774     for (struct proc *p1 = ptable.proc; p1 < &ptable.proc[NPROC]; p1++){
775         if(p1->sched_policy == 1 || p1->pid == p->pid){
776             //calculating utilization of all sched_policy=1 including the new process
777             num_rms_proc++;
778             cpu_utilz += (float)(p1->rate*p1->exec_time);
779         }
780     }
781
782     if (cpu_utilz > rms_bound[num_rms_proc-1]){
783         release(&ptable.lock);
784         kill(pid);
785         return -22;
786     }
787 }
```

Scheduling Algorithm

In RMS scheduling, processes are prioritized based on their rates or priorities, with higher priority processes scheduled first. Our scheduling algorithm calculates the

weight of each process and schedules the one with the least weight. In cases where multiple processes have the same weight, the process with the lower PID is given precedence. Here is a simplified version of the scheduling algorithm code:

```
427 else if (sch_policy == 1) //RMS
428 {
429
430     int min_priority = 5;
431     struct proc *p_selected = NULL;
432
433     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
434     {
435         if (p->state == RUNNABLE && p->sched_policy == 1){
436             int priority = rms_priority(p);
437
438             if (priority < min_priority){ //find minimum priority
439                 min_priority = priority;
440                 p_selected=p;
441             }
442             else if (priority == min_priority){ //select lower pid
443                 if (p->pid < p_selected->pid)
444                     p_selected=p;
445             }
446         }
447     }
448
449     if (p_selected != NULL){
450         p = p_selected;
451         // Switch to chosen process. It is the process's job
452         // to release ptable.lock and then reacquire it
453         // before jumping back to us.
454         c->proc = p;
455         switchvm(p);
456         p->state = RUNNING;
457
458         swtch(&(c->scheduler), p->context);
459         switchkvm();
460
461         p->elapsed_time++;
462         // Process is done running for now.
463         // It should have changed its p->state before coming back.
464         c->proc = 0;
465     }
466 }
```

The implementation of the RMS scheduling algorithm in xv6 enhances the real-time capabilities of the operating system by efficiently scheduling processes based on their rates or priorities. The schedulability test ensures that processes are admitted into the scheduling queue appropriately, while the scheduling algorithm prioritizes processes effectively to meet real-time requirements. Overall, the RMS algorithm contributes to the reliability and responsiveness of the xv6 operating system in handling real-time tasks.

Modifications in the proc Structure:

- `sched_policy`: This attribute denotes the scheduling policy of the process. It can take values such as -1 for the default xv6 policy, 0 for EDF (Earliest Deadline First), 1 for RMS (Rate-Monotonic Scheduling), and 4 for non-schedulable processes. Users can set this attribute using the `sched_policy(pid, value)` system call.
 - `elapsed_time`: It tracks the elapsed time of the process by counting the number of ticks during which the process was in the RUNNING state.
 - `exec_time`: This attribute specifies the total allowed execution time for the process, settable by user processes using the `exec_time(pid, value)` system call.
 - `rate`: Represents the rate of assumed periodic processes. Users can set this attribute using the `rate(pid, value)` system call.
 - `deadline`: Denotes the hard deadline of the process. If a new process fails the schedulability check, it is killed to ensure that deadlines for accepted processes are not compromised. The deadline is set using the `deadline(pid, value)` system call.
 - `arrival_time`: Represents the start time of the process, recording the tick value when the process's `sched_policy` was set.
-

```

37 // Per-process state
38 struct proc {
39     uint sz;                // Size of process memory (bytes)
40     pde_t* pgdir;          // Page table
41     char *kstack;          // Bottom of kernel stack for this process
42     enum procstate state;   // Process state
43     int pid;               // Process ID
44     struct proc *parent;    // Parent process
45     struct trapframe *tf;   // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan;            // If non-zero, sleeping on chan
48     int killed;            // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;     // Current directory
51     char name[16];         // Process name (debugging)
52
53     int deadline;          // deadline of the process
54     int exec_time;         // execution time of the process
55     int sched_policy;       // scheduling policy of the process, -1 for round robin 0 for edf, 1 for rms
56     int elapsed_time;      // elapsed time of the process
57     int rate;              // rate of the process
58     int arrival_time;      // arrival time of the process
59
60 };
61

```

trap.c Modifications

The trap.c file has been modified to handle unwanted or completed processes efficiently. A process is considered unwanted if its sched_policy is not -1 (indicating a specific scheduling policy) and its elapsed_time exceeds its exec_time. The modified code snippet below illustrates this functionality:

```

103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER){
106
107     if((myproc()->sched_policy != 0) && (myproc()->elapsed_time >= myproc()->exec_time)){
108         printf("The completed process has arrival time and pid values: %d %d\n", myproc()->arrival_time, myproc()->pid);
109         exit();
110     }
111     else{
112         yield();
113     }
114 }
115

```

RMS priority

This code calculates priorities of rate monotonic scheduled processes based on rate values.

```

13 int rms_priority(struct proc *p){
14     int numerator = (30 - p->rate)*3;
15     int priority = numerator / 29;
16     if(priority<1){
17         priority=1;
18     }
19     else if(numerator % 29 != 0){
20         priority += 1;
21     }
22     return priority;
23 }

```
